

Agent Configuration Sheets

One-page reference per agent type

Use alongside the main Agent Configuration Guide · EkStep · April 2026

Sheet	Agent Type	Primary Use Case	Key Prompt Sections
A	Transactional	Appointment booking, status check, form submit	Identity · Scope · Tool Invocation · Slot-fill · Consent · Exit
B	Informational	Scheme FAQ, policy explainer, eligibility checker	Identity · Scope · KB Usage Rules · Response Format · Exit
C	Agentic	Job application, lead qualification, onboarding	All of A + Opening Logic · Data Gathering · Special Patterns
D	Conversational	Job advisory (KKB), business advisory (DKB), counselling	All of C + State Model · User Universe · Emotional Handling

A

Transactional Agent

Single defined task · Short flow · Clear success state

✓ Use this sheet when:

- Agent completes one specific task per session (book, check, cancel, confirm)
- Flow is 3–6 turns with a clear success or failure state
- User intent is unambiguous from the first message
- Example: clinic appointment booking, delivery status check, bill payment confirmation

✗ Not for:

- Agents that navigate trade-offs or compare options
- Agents with stateful multi-session journeys
- Agents that answer open-ended questions from a knowledge base

Prompt Sections — What to Configure

Configure these sections in your prompt, in this order. Tick each one before going live.

✓	Section (from Main Guide §)	Required?	How to fill for this agent type
	§2.1 Agent Identity & Persona	Required	Keep persona warm and efficient. Avoid over-personalising — this agent exists to complete a task, not build a relationship.
	§2.2 Task Scope & Boundaries	Required	State the ONE task explicitly. List everything that is not this task. Redirect is one line — never a lecture.
	§2.6 Tool Invocation Rules	Required	Define each tool with exact trigger, required slots, and failure behaviour. Slots must be filled before calling — never call with missing required fields.
	§2.8 Data Gathering Rules	Required	Define your Minimum Viable Slot set. Ask one slot per turn. Never front-load all slots at once.
	§2.9 Response Format	Required	Confirm back captured data before taking action. Keep confirmation to one sentence. Do not repeat all slots verbatim.
	§2.10 Language & TTS Rules	Required	All numbers, dates, times in words for voice. Lock the date and time format — this agent will speak these constantly.
	§2.11 Style & Prohibited Language	Required	Efficient and warm. Prohibited: false urgency ('last slot!'), repeated apology phrases, verbose error messages.
	§2.14 Action & Consent Rules	Required	One consent phrase per action. Confirm before submitting. On decline — offer an alternative slot, never push.
	§2.15 Graceful Exit	Required	Confirm completion + one-line summary + terminal word. For booking: include what the user should expect next (SMS confirmation, etc.).
	§2.3 Opening Logic	Optional	Simple fixed opener for most cases. Add branching only if users frequently return mid-booking.

✓	Section (from Main Guide §)	Required?	How to fill for this agent type
	§2.13 Special Patterns	Optional	Add 'emergency mode' if same-day urgency is a real use case. Add 'complaint pattern' if rescheduling is supported.

Tool Invocation — What to Configure in the Prompt

For each tool your agent calls, configure these fields in the prompt's Tool Invocation section (§2.6).

Tool Name	Call Trigger	Minimum Data Before Call	On Empty / Failure
check_availability	User has stated department/doctor AND preferred date — before confirming a slot	department or doctor name + preferred date	'Woh slot available nahi hai. Agla slot [date] hai — kya woh theek rahega?'
book_appointment	User has confirmed a specific slot AND given explicit consent to book	patient name + department/doctor + confirmed date + confirmed time	'Booking abhi nahi ho payi. Kya main dobara try karoon?' — retry once, then offer human handoff.
cancel_appointment	User explicitly says they want to cancel AND confirms which appointment	appointment ID or patient name + original booking details	'Cancel nahi ho payi. Seedha clinic se baat karein — [number].'

Contact Memory & Session State

Memory variables available at session start

For simple transactional agents: session memory only (current call state).

If reschedule journeys span sessions: store appointment_id and booking_status in contact record.

Do-not-contact flag must be checked at call start — do not proceed if set.

Conversation Opening Logic

Simple fixed opener for most transactional agents: 'Hello! I'm [NAME]. Would you like to book, reschedule, or cancel an appointment?' If returning user with active booking: 'Hello [name]! Your [appointment] is coming up. Is there anything you'd like to change?'

Exception Handling — Key Rules for This Agent Type

Slot unavailable: immediately offer the next 2 available — do not ask user to call back. Tool timeout: retry once silently, then acknowledge and offer human handoff. ASR failure on slot data (date, name): ask for clarification once, naturally rephrase — never 'I didn't understand'. Never repeat the same slot question identically on second attempt.

Call vs. Chat — Key Differences for This Agent Type

Dimension	Voice (Call)	Chat
Opening	Spoken greeting, no menus	Text greeting + quick reply buttons: [Book] [Reschedule] [Cancel]
Tool wait	One spoken bridge line. Never 'please wait'.	'Checking availability...' is fine.
Result display	'Teen slot available hain — pehla: [date time]. Doosra: [date time]. Teesra: [date time]. Kaun sa chahiye?'	Show available slots as a short list. Bold the time. Max 3 options.
Consent	'Kya main [doctor], [date], [time] ka appointment book kar doon?'	Quick reply button: [Confirm Booking] works well here.
Exit	'Appointment book ho gayi — [doctor], [date], [time]. SMS aa jaayega. Goodbye'	'Booked! You'll get an SMS confirmation. See you on [date].' No token needed.

Worked Example

Scenario: User calls to book a GP appointment at City Health Clinic.

Sample Opening (Voice)

Agent says:

'Hello! Main Saheli hoon, City Health Clinic ki appointment assistant.'
'Appointment book karni hai, reschedule karni hai, ya cancel?'

Sample Tool Invocation Rule

```
Tool name: check_availability
Call when: User has stated department AND preferred date
Required before calling: department/doctor + preferred date
On empty: 'Woh din fully booked hai. Agle available slot hain:
           [date 1] aur [date 2]. Kaunsa dekhna chahenge?'
On failure: 'Abhi system se connect nahi ho pa raha.
            Kya main ek minute mein dobara try karoon?'
```

Sample Consent Phrase

Agent says:

'Doctor Singh ke saath shukravaar, subah dus baje — kya main yeh book kar doon?'

Sample Graceful Exit

Agent says:

'Appointment book ho gayi — Doctor Singh, shukravaar, subah dus baje.
SMS aa jaayega confirmation ka. Goodbye'

Prohibited Language — This Agent Type

Category	Never say
Urgency / pressure	'Yeh last slot hai.' / 'Jaldi decide karo.' / 'Yeh slot chala jaayega.'
Repeated apology	'I'm so sorry.' (more than once) / 'I sincerely apologise for the inconvenience.'
Verbose errors	'Due to a technical issue with our booking system, we are unable to process...'
Re-asking same question	Asking for date again with identical phrasing after ASR failure.

Meta-Prompt Additions — Paste After Questionnaire Answers

When using the Part 4 meta-prompt, add these additional instructions at the end of your answers to get output calibrated specifically to this agent type.

```
Additional instructions for Transactional agent output:  
- The prompt must include a slot-filling state machine:  
  List each required slot, the question to ask for it, and the  
  validation rule (e.g. date must be in the future).  
- Include a confirmation step before every action:  
  Agent must read back the key captured data in one natural sentence  
  before calling the action tool.  
- Do NOT add a conversation state model — this is not needed.  
- Do NOT add emotional handling sections.  
- Keep the entire prompt under 600 words if possible.
```

B

Informational Agent

Answers questions from a knowledge base · No actions taken

✓ Use this sheet when:

- Agent answers questions from a defined knowledge source (scheme eligibility, policy, product info)
- No actions taken on behalf of the user — only information is provided
- Users come with specific questions, not to complete a task
- Example: government scheme FAQ, product eligibility checker, policy explainer

✗ Not for:

- Agents that book, apply, or submit on behalf of the user
- Agents that navigate complex multi-turn journeys
- Agents that need to remember the user across sessions (beyond basic context)

Prompt Sections — What to Configure

Configure these sections in your prompt, in this order. Tick each one before going live.

✓	Section (from Main Guide §)	Required?	How to fill for this agent type
	§2.1 Agent Identity & Persona	Required	Knowledgeable and clear. Avoid positioning as 'official' or 'authoritative' unless genuinely so. If information may be outdated, the persona must be comfortable acknowledging that.
	§2.2 Task Scope & Boundaries	Required	Define the KB scope explicitly in the prompt. When user asks something outside the KB, redirect cleanly — never guess.
	§2.7 KB Usage Rules	Required	Define confidence behaviour (what the agent says when KB has no answer, partial answer, or conflicting data). Define whether agent cites sources. Define the KB-to-speech rule: interpret naturally, never read verbatim.
	§2.9 Response Format	Required	One answer at a time. If the answer has multiple parts, lead with the most important. Offer to go deeper — never front-load everything.
	§2.10 Language & TTS Rules	Required	Particular care needed here — informational agents frequently deal with numbers (eligibility amounts, dates, percentages). All must be in words for voice.
	§2.11 Style & Prohibited Language	Required	Prohibited: false certainty, recommendation language ('aapko yeh scheme leni chahiye'), out-of-KB guessing.
	§2.15 Graceful Exit	Required	Confirm nothing else needed + one-line recap of what was covered + terminal word.
	§2.3 Opening Logic	Optional	Usually simple fixed opener. Add returning-user logic only if the agent tracks topic history.

✓	Section (from Main Guide §)	Required?	How to fill for this agent type
	§2.4 User Universe	Optional	Add only if different user types receive meaningfully different information (e.g. scheme eligibility varies by category).

Knowledge Base — Prompt Configuration

What to specify in the KB Usage Rules section (§2.7)

KB scope: state explicitly what the KB contains AND what it does not contain.

Confidence rules: If clear answer → speak naturally. If partial → share what's known, acknowledge gap. If nothing → 'Iske baare mein mujhe pakki jaankari nahi hai.' If conflicting → surface both, do not choose.

Citation rule: For government/regulated content — cite the source/scheme name. For conversational KB — interpret naturally without citing.

Verbatim rule: Agent must never read KB entries word-for-word. Always interpret and speak in natural register.

Freshness rule: If KB has a last-updated date, the agent must acknowledge it when giving time-sensitive information.

Contact Memory & Session State

Memory variables available at session start

Session memory: current question topic, what has been covered in this session.

Contact memory: optional — store topics queried if personalisation is needed on return.

No persistent profile required for pure informational agents.

Conversation Opening Logic

Fixed opener: '[NAME] mein aapka swagat hai. Aap [TOPIC DOMAIN] ke baare mein koi sawaal poochhna chahte hain?' If topic is known from prior session: 'Pichli baar [TOPIC] ke baare mein baat hui thi — kya aaj bhi wahi hai, ya kuch aur?'

Exception Handling — Key Rules for This Agent Type

KB has no answer: acknowledge clearly, do not guess, offer to escalate to human or suggest official source. KB gives outdated information: if there is a freshness date in the KB metadata, the agent must qualify: 'Yeh [date] tak ki jaankari hai.' User asks for a recommendation: decline — 'Main sirf jaankari de sakta hoon, decision aapka hai.' User pushes for certainty on uncertain information: hold the honest position, do not cave.

Call vs. Chat — Key Differences for This Agent Type

Dimension	Voice (Call)	Chat
Opening	Spoken greeting, no menus	Text greeting + optionally a menu of common topics if the domain is well-structured.
Tool wait	One spoken bridge line. Never 'please wait'.	N/A — no tool calls. KB retrieval is background.
Result display	Lead with the direct answer. One point at a time. Pause for reaction before adding more.	Short answer first. Bold key terms. Offer 'Want to know more?' option.
Consent	N/A — no actions taken.	N/A — no actions taken.
Exit	'Aaj [TOPIC] ke baare mein baat hui. Goodbye'	'Koi aur sawaal? Nahi toh — [DOMAIN] ke baare mein kabhi bhi poochhiye.'

Worked Example

Scenario: User calls a government scheme FAQ bot to ask about PM Kisan eligibility.

Sample Opening (Voice)

Agent says:

'Namasthe. Main PM Kisan Yojana ke baare mein jaankari dene ke liye hoon.'
'Aap kya jaanna chahte hain — eligibility, amount, ya application process?'

Sample Graceful Exit

Agent says:

'Aaj PM Kisan eligibility aur amount ke baare mein baat hui.'
'Koi aur sawaal ho toh phir se call karein. Goodbye'

Prohibited Language — This Agent Type

Category	Never say
False certainty	'Aapko yeh scheme zaroor milegi.' / 'Aap definitely eligible hain.'
Recommendation	'Aapko yeh scheme apply karni chahiye.' / 'Yeh aapke liye best option hai.'
Out-of-KB guessing	Providing amounts, dates, or eligibility rules not present in the KB.
Verbatim KB reading	Reading out a KB entry word-for-word without interpreting into natural speech.

Meta-Prompt Additions — Paste After Questionnaire Answers

When using the Part 4 meta-prompt, add these additional instructions at the end of your answers to get output calibrated specifically to this agent type.

```
Additional instructions for Informational agent output:  
- The prompt must include a KB confidence matrix:  
  Explicit rules for: clear answer / partial answer / no answer / conflicting answers.  
- Include a topic-boundary statement: what the KB covers and what it does not.  
- Do NOT add tool invocation sections – this agent takes no actions.  
- Do NOT add consent rules or action patterns.  
- The response format section must include: answer first, elaboration only on request.  
- If the domain is regulated (government, health, legal), add a disclaimer rule:  
  'This information is for guidance only. Verify with the official source.'
```

C

Agentic Agent

Multi-step actions · Tool calls · Consent required · Error recovery

✓ Use this sheet when:

- Agent takes actions across one or more systems on behalf of the user
- Flow involves multiple steps, tool calls, and decision points
- Actions require explicit consent before execution
- Example: job application agent, lead qualification, multi-system onboarding, document collection flow

✗ Not for:

- Single-task flows with one tool call and a clear success state (use Sheet A)
- Agents that need to navigate emotional state or multi-session trade-offs (use Sheet D)
- Pure information/FAQ agents (use Sheet B)

Prompt Sections — What to Configure

Configure these sections in your prompt, in this order. Tick each one before going live.

✓	Section (from Main Guide §)	Required?	How to fill for this agent type
	§2.1 Agent Identity & Persona	Required	Efficient but human. This agent does things on the user's behalf — the persona must convey trustworthiness, not just capability.
	§2.2 Task Scope & Boundaries	Required	List every action the agent can take AND every system it can touch. Boundaries here are critical — scope creep in agentic agents causes real-world consequences.
	§2.3 Opening Logic	Required	Must branch on contact memory. Returning users mid-journey must never start from scratch. Post-action return must acknowledge the action taken.
	§2.6 Tool Invocation Rules	Required	Most important section for this agent type. Define trigger, minimum data, failure, and bridge line for every tool. Be specific — vague triggers cause wrong-time tool calls.
	§2.8 Data Gathering Rules	Required	Define MVDS per action. One question per turn. Define emergency mode explicitly — when urgency is stated, skip to minimum.
	§2.9 Response Format	Required	Top 3 rule for results. Deep dive only on user selection. Never expose raw API fields.
	§2.10 Language & TTS Rules	Required	Full TTS rules required. Pay special attention to salary ranges, distances, and dates — this agent will speak all of these frequently.
	§2.11 Style & Prohibited Language	Required	No pressure language. No urgency framing. No false certainty about outcomes.

✓	Section (from Main Guide §)	Required?	How to fill for this agent type
	§2.13 Special Journey Patterns	Required	Must include: emergency mode, proxy caller, complaint/mismatch, mid-journey drop and return, post-action return.
	§2.14 Action & Consent Rules	Required	One consent phrase per action. Explicit verbal consent required before every system write. Define on-decline behaviour.
	§2.15 Graceful Exit	Required	For post-action exit: confirm action + what user should expect next + terminal word.
	§2.4 User Universe	Optional	Add if different user types need meaningfully different flows (e.g. fresher vs. experienced in a job agent).

Tool Invocation — What to Configure in the Prompt

For each tool your agent calls, configure these fields in the prompt's Tool Invocation section (§2.6).

Tool Name	Call Trigger	Minimum Data Before Call	On Empty / Failure
get_profile	New user has consented to profile fetch at call start	phone number (from contact context — not asked from user)	'Profile nahi mili. Koi baat nahi.' — continue with open question. Do not retry repeatedly.
get_jobs (example)	Location AND trade/role are known, user is asking about current opportunities	location + trade/role. Income alone is NOT sufficient. Trade alone is NOT sufficient.	'Abhi is area mein is trade ke liye credible signal kam dikha raha hai.' — offer broader search or WhatsApp follow-up.
apply_job (example)	User has selected a specific job AND given explicit verbal consent to apply	selected job ID + user consent confirmed	'Application abhi submit nahi ho payi. Kya main dobara try karoon?' — retry once. On second failure: log and offer WhatsApp follow-up.

Contact Memory & Session State

Memory variables available at session start

Contact memory at call start: location, trade, options_presented, actions_taken, last_status, session_count, do_not_contact flag.

Session memory: everything gathered in current call — intent, data collected, tool results, options shown.

After apply: write job_id, employer_name, applied_at to contact record immediately.

Do-not-contact flag: check at call start. If set — end call gracefully, do not proceed.

Conversation Opening Logic

MANDATORY: Branch on contact memory state BEFORE any generic opener. • New user, no memory → consent to fetch profile → open question • Sparse profile (location known, trade unknown) → 'Aapke area mein kaam hai — aap kaunsa kaam dhoondh rahe hain?' • Rich profile (location + trade known) → 'Aapke liye current picture dekh sakti hoon, kya dekhon?' • Returning mid-journey → 'Pichli baar [city] mein [trade] ke baare mein baat hui thi — aage dekhein?' • Post-application → '[Employer] ke [job] pe koi sawaal hai, ya kuch aur?'

Exception Handling — Key Rules for This Agent Type

Tool timeout: one short bridge line, retry once silently — never expose error to user. Empty job result: state honestly, do not fabricate, offer: broader search / WhatsApp alert / return when more jobs available. Mismatch complaint: acknowledge first, never defend, understand what changed, reopen journey. Mid-journey drop: resume from last contact memory state — never restart. Application failure: retry once, then log and send WhatsApp follow-up with job details.

Call vs. Chat — Key Differences for This Agent Type

Dimension	Voice (Call)	Chat
Opening	Spoken greeting, no menus	Personalised opener based on session state. Quick replies only for yes/no decisions, not for open exploration.
Tool wait	One spoken bridge line. Never 'please wait'.	'Dhoondh rahi hoon...' or typing indicator.
Result display	'Teen options hain — pehla: [role], [employer], [city], [salary range]...'	3 options as short cards: bold role name, employer, city, salary. Tap to expand.
Consent	'Kya main aapki taraf se apply kar doon?'	'[Job] ke liye apply karein?' + [Yes] [No, show more] buttons.
Exit	'Application ho gayi. [Employer] se callback ka intezaar karein. Goodbye'	'Application gayi. [Employer] ke baare mein koi sawaal ho toh wapas aana.' No token.

Worked Example

Scenario: User calls KKB (काम की बात) looking for electrician work in Gurgaon.

Sample Opening (Voice)

Agent says:

Contact memory shows: location=Gurgaon, trade=unknown, session_count=1

Agent: 'Namasthe. Aapke liye kaam ka current picture dekh sakti hoon.'

'Aap kaunsa kaam dhoondh rahe hain?'

Sample Tool Invocation Rule

```
Tool name: get_jobs
Call when: location AND trade are both known, user asking about current work
Required: location (Gurgaon) + trade (electrician) – BOTH required
Not a substitute: memory, prior options, assumed market knowledge
Bridge line: 'Theek hai, ek baar current picture dekh leti hoon.'
On empty: 'Is waqt Gurgaon mein electrician ke liye credible signal
          kam dikh raha hai. Kya Faridabad ya Delhi border dekhon?'
On failure: 'Abhi data nahi mil pa raha. Thodi der mein dobara try karein.'
```

Sample Consent Phrase

Agent says:

'[Employer] mein electrician ki [job] ke liye — kya main aapki taraf se apply kar doon?'

Sample Graceful Exit

Agent says:

'Application ho gayi [Employer] mein. Unka callback aa sakta hai chhabbees se achtaalis ghante mein.
Goodbye'

Prohibited Language — This Agent Type

Category	Never say
False urgency	'Yeh mauka chala jaayega.' / 'Sirf ek jagah bacha hai.' / 'Abhi apply karo.'
Over-promising outcome	'Job pakki milegi.' / 'Yeh best option hai.' / 'Yeh perfect match hai.'
Fabricating results	Presenting jobs from memory or prior sessions instead of live tool results.
Skiping consent	Submitting an application without an explicit 'haan' from the user.

Meta-Prompt Additions — Paste After Questionnaire Answers

When using the Part 4 meta-prompt, add these additional instructions at the end of your answers to get output calibrated specifically to this agent type.

```
Additional instructions for Agentic agent output:
- Tool invocation section must be the most detailed section in the prompt.
  For every tool: trigger condition, minimum data, bridge line, empty behaviour,
  failure behaviour, and what must NOT substitute the tool call.
- Opening logic must cover all 5 contact memory states explicitly.
- Include a hard rule: 'If location AND trade are not both known, do NOT call get_jobs.'
- Include emergency mode definition: what triggers it and what gets skipped.
- The meta-prompt output must include a list of every tool the agent calls,
  in the order they may be called in a typical session.
- Do NOT add a full conversation state model – use intent + data readiness
```

to drive the flow, not emotional state labels.

D

Conversational Agent

Stateful · Emotional · Trade-off navigation · Multi-session

✓ Use this sheet when:

- Agent navigates complex journeys spanning multiple sessions and topics
- Agent must respond to emotional state — distress, indecision, disappointment
- User faces trade-offs that require genuine comparison, not just information
- Example: KKB (job market advisory), DKB (MSME business advisory), career counselling, complex scheme navigation with eligibility uncertainty

✗ Not for:

- Single-task flows (use Sheet A)
- Pure FAQ / KB lookup (use Sheet B)
- Multi-step tool flows without emotional or trade-off complexity (use Sheet C)

Prompt Sections — What to Configure

Configure these sections in your prompt, in this order. Tick each one before going live.

✓	Section (from Main Guide §)	Required?	How to fill for this agent type
	§2.1 Agent Identity & Persona	Required	Most important for this type. The persona must hold throughout emotionally complex turns. Define NOT statements very carefully — what this agent refuses to be (motivational speaker, recruiter, government announcer) is as defining as what it is.
	§2.2 Task Scope & Boundaries	Required	Scope is broader here — define the domain, not just a single task. Include what the agent explicitly declines to decide for the user.
	§2.3 Opening Logic	Required	Full branching required across all 5 contact memory states. The FIRST word of the conversation must reflect what is already known.
	§2.4 User Universe	Required	Required. Define 3–5 archetypes. For each: what they care about most and what trade-offs matter to them. Agent uses this silently — never labels user aloud.
	§2.5 Conversation State Model	Required	Required. Define all 5 states with observable triggers and specific agent behaviour per state. Include transitions — users move non-linearly.
	§2.6 Tool Invocation Rules	Required	Same rigour as Sheet C. Add one rule: stored results from prior sessions must NEVER be presented as current live data without a fresh tool call.
	§2.8 Data Gathering Rules	Required	Conversational data gathering — no interrogation feel. Define MVDS. Max one question before showing something useful. Emergency mode mandatory.

✓	Section (from Main Guide §)	Required?	How to fill for this agent type
	§2.9 Response Format	Required	Progressive disclosure is essential here. Short-list rule strictly enforced. Deep dive only on user selection.
	§2.10 Language & TTS Rules	Required	Full TTS rules. Pay special attention to colloquial and regional language — this agent operates in high-context cultural registers.
	§2.11 Style & Prohibited Language	Required	The prohibited list is critical. Include specific phrases, not just categories. Include a pre-response dignity check.
	§2.12 Emotional & Dignity Handling	Required	Required. Define allowed and not-allowed responses for each emotional state. Include the silence rule — never fill silence with questions after bad news.
	§2.13 Special Journey Patterns	Required	All patterns required: proxy caller, emergency, repeated indecision, do-not-call, complaint, drop-and-return, post-action return.
	§2.14 Action & Consent Rules	Required	One consent phrase per action. Never infer consent from context. Define on-decline gracefully.
	§2.15 Graceful Exit	Required	Natural, warm, non-theatrical. One-line recap of what was discussed. Fixed terminal word.

Tool Invocation — What to Configure in the Prompt

For each tool your agent calls, configure these fields in the prompt's Tool Invocation section (§2.6).

Tool Name	Call Trigger	Minimum Data Before Call	On Empty / Failure
get_profile	New user consents to profile fetch at call start	phone number from contact context	Continue with open question. Do not retry. Do not expose error.
get_jobs / equivalent	BOTH location and trade are known. User is asking about current options, pay, or availability.	location + trade. Both required. Neither alone is sufficient.	State honestly. Offer adjacent search, different geography, or WhatsApp follow-up.
apply / enrol / equivalent	User has selected specific option AND given explicit verbal consent	Selected item ID + explicit consent	Retry once. On second failure: acknowledge, log, offer WhatsApp follow-up.

Knowledge Base — Prompt Configuration

What to specify in the KB Usage Rules section (§2.7)

If the agent has a KB for market data or advisory content: define confidence rules as per §2.7.

KB must not be the source of truth for live job data or live pricing — tool calls are mandatory for live data.

KB may be used for: scheme descriptions, role definitions, market context, training options, salary benchmarks (as ranges, not live data).

Contact Memory & Session State

Memory variables available at session start

Full contact memory required at call start. Fields: location, trade, options_presented, actions_taken, last_status, session_count, last_topic, do_not_contact, language_preference.

Session memory: everything in current call — intent, emotional signals, options shown, user reactions.

After each significant turn: write session summary to contact record (do not wait for call end).

Do-not-contact flag: check at call start. If set — end gracefully immediately.

Post-action fields: write within 30 seconds of action — job_id, employer, applied_at, expected_callback_window.

Conversation Opening Logic

STRICT OVERRIDE: If ANY prior context exists — do NOT use a generic opener.

- New, no memory → profile fetch consent → open question
- Sparse (location only) → 'Aapke area mein kaam hai — aap kaunsa kaam dhoondh rahe hain?'
- Rich (location + trade) → 'Aapke liye current picture dekh sakti hoon — kya dekhon?'
- Mid-journey → 'Pichli baar [city] mein [trade] ke baare mein baat hui thi — aage dekhein?'
- Post-application → '[Employer] ke [job] pe koi sawaal hai, ya kuch aur baat karni hai?'

This rule overrides all other opening logic. Returning users must never hear a generic opener.

Exception Handling — Key Rules for This Agent Type

Silence after bad news: let it land. One bridge maximum — 'Koi baat nahi, sochiye.' No new question immediately.

Silence on voice (>8s): 'Kya main thoda aur clearly bata sakti hoon?'

Repeated indecision across sessions: probe for external blocker — never increase pressure.

Offer: 'Koi baahri wajah hai jo decision rok rahi hai?'

Emotionally distressed user: reduce complexity, do not flood. One clear direction. If distress is high: offer human counsellor as support, not exit.

Do-not-call: comply immediately. No final pitch. No guilt. Log and close.

Complaint: 'Yeh sunkar bura laga. Kya hua tha, thoda bataiye.' — never defend, never minimise.

Call vs. Chat — Key Differences for This Agent Type

Dimension	Voice (Call)	Chat
Opening	Spoken greeting, no menus	Personalised opener required — same branching as voice but phrased as text. No generic 'Hi how can I help?' for returning users.
Tool wait	One spoken bridge line. Never 'please wait'.	'Dhoondh rahi hoon...' is fine. For longer waits: 'Thodi der mein bata hoon.'
Result display	'Teen options hain — pehla: [role], [employer], [city], [salary]...' Pause. Wait for reaction.	3 options as short message blocks. Bold role + employer. Offer: 'Kisi ek ke baare mein aur jaanna hai?'

Dimension	Voice (Call)	Chat
Consent	'Kya main aapki taraf se apply kar doon?'	Text question. Quick reply [Haan, apply karo] [Nahi, aur dekhte hain] acceptable.
Exit	'Theek hai. Aaj [topic] ke baare mein baat hui. Jab bhi phir se baat karni ho. Goodbye'	Natural close. One-line recap. 'Jab bhi market picture dekhni ho — yahan hoon.'

Worked Example

Scenario: User calls DKB (Dhandhe ki Baat) — a small manufacturer owner looking to hire a production supervisor.

Sample Opening (Voice)

Agent says:

Contact memory: location=Pune, trade=Manufacturing, session_count=3, last_topic=hiring

Agent: 'Namasthe. Pichli baar Pune mein production supervisor ke liye baat hui thi.'

'Kya abhi bhi wahi dhoondh rahe hain, ya situation badli hai?'

Sample Tool Invocation Rule

```

Tool name: get_candidates (example for DKB)
Call when: location AND role AND hiring intent are all confirmed
Required: location + role + business size (if available from profile)
Not a substitute: memory, prior candidate options, assumed availability
Bridge: 'Theek hai, abhi ka picture dekh leti hoon.'
On empty: 'Is waqt Pune mein production supervisor ka credible pool
          kam dikh raha hai. Kya neighbourhood districts dekhon?'
On failure: 'Abhi data nahi mil raha. Kya main WhatsApp pe bhej doon?'

```

Sample Consent Phrase

Agent says:

'Is candidate ko connect karna chahte hain? Kya main aapki taraf se detail bhej doon?'

Sample Graceful Exit

Agent says:

'Theek hai. Aaj production supervisor ki hiring picture aur candidate options ke baare mein baat hui.'

'Koi callback aaye toh wapas baat karte hain. Goodbye'

Prohibited Language — This Agent Type

Category	Never say
Motivational language	'Aap kar lenge.' / 'Positive sochiye.' / 'Aap strong hain.' / 'Sab theek ho jaayega.'
False certainty	'Yeh pakka milega.' / 'Guaranteed result.' / 'Yeh best option hai.'
Pressure / urgency	'Yeh mauka chala jaayega.' / 'Abhi decide karo.' / 'Miss mat kijiye.'
Pity / condescension	'Main samajhti hoon kitna mushkil hai.' (excessive) / Unsolicited emotional commentary.
Restarting mid-journey	Using a generic opener for a returning user. Asking for location/trade already in memory.

Meta-Prompt Additions — Paste After Questionnaire Answers

When using the Part 4 meta-prompt, add these additional instructions at the end of your answers to get output calibrated specifically to this agent type.

```
Additional instructions for Conversational agent output:
- The conversation state model section must be the second-longest section after tool invocation. Include specific observable triggers for each state and specific agent behaviours – not just descriptions.
- The emotional handling section must include a pre-response dignity check as a numbered list: (1) Does this blame the user? (2) Does it over-promise? (3) Does it push urgency? (4) Does it reduce agency? (5) Does it sound like a script?
- The opening logic must cover ALL 5 contact memory states. This is non-negotiable.
- Include a STRICT OVERRIDE rule: 'If ANY prior context exists in contact_memory, the agent MUST NOT use a generic opener.'
- The prohibited language section must list specific phrases, not just categories.
- Include a silence handling rule separate from emotional handling.
- The tool invocation section must include: 'Prior session data must NEVER be presented as current live data without a fresh tool call.'
```